MUSTER

Strong retrieval without generative construction

BEAM-10M / GLM-5.3-Flash reader / official rubric judge

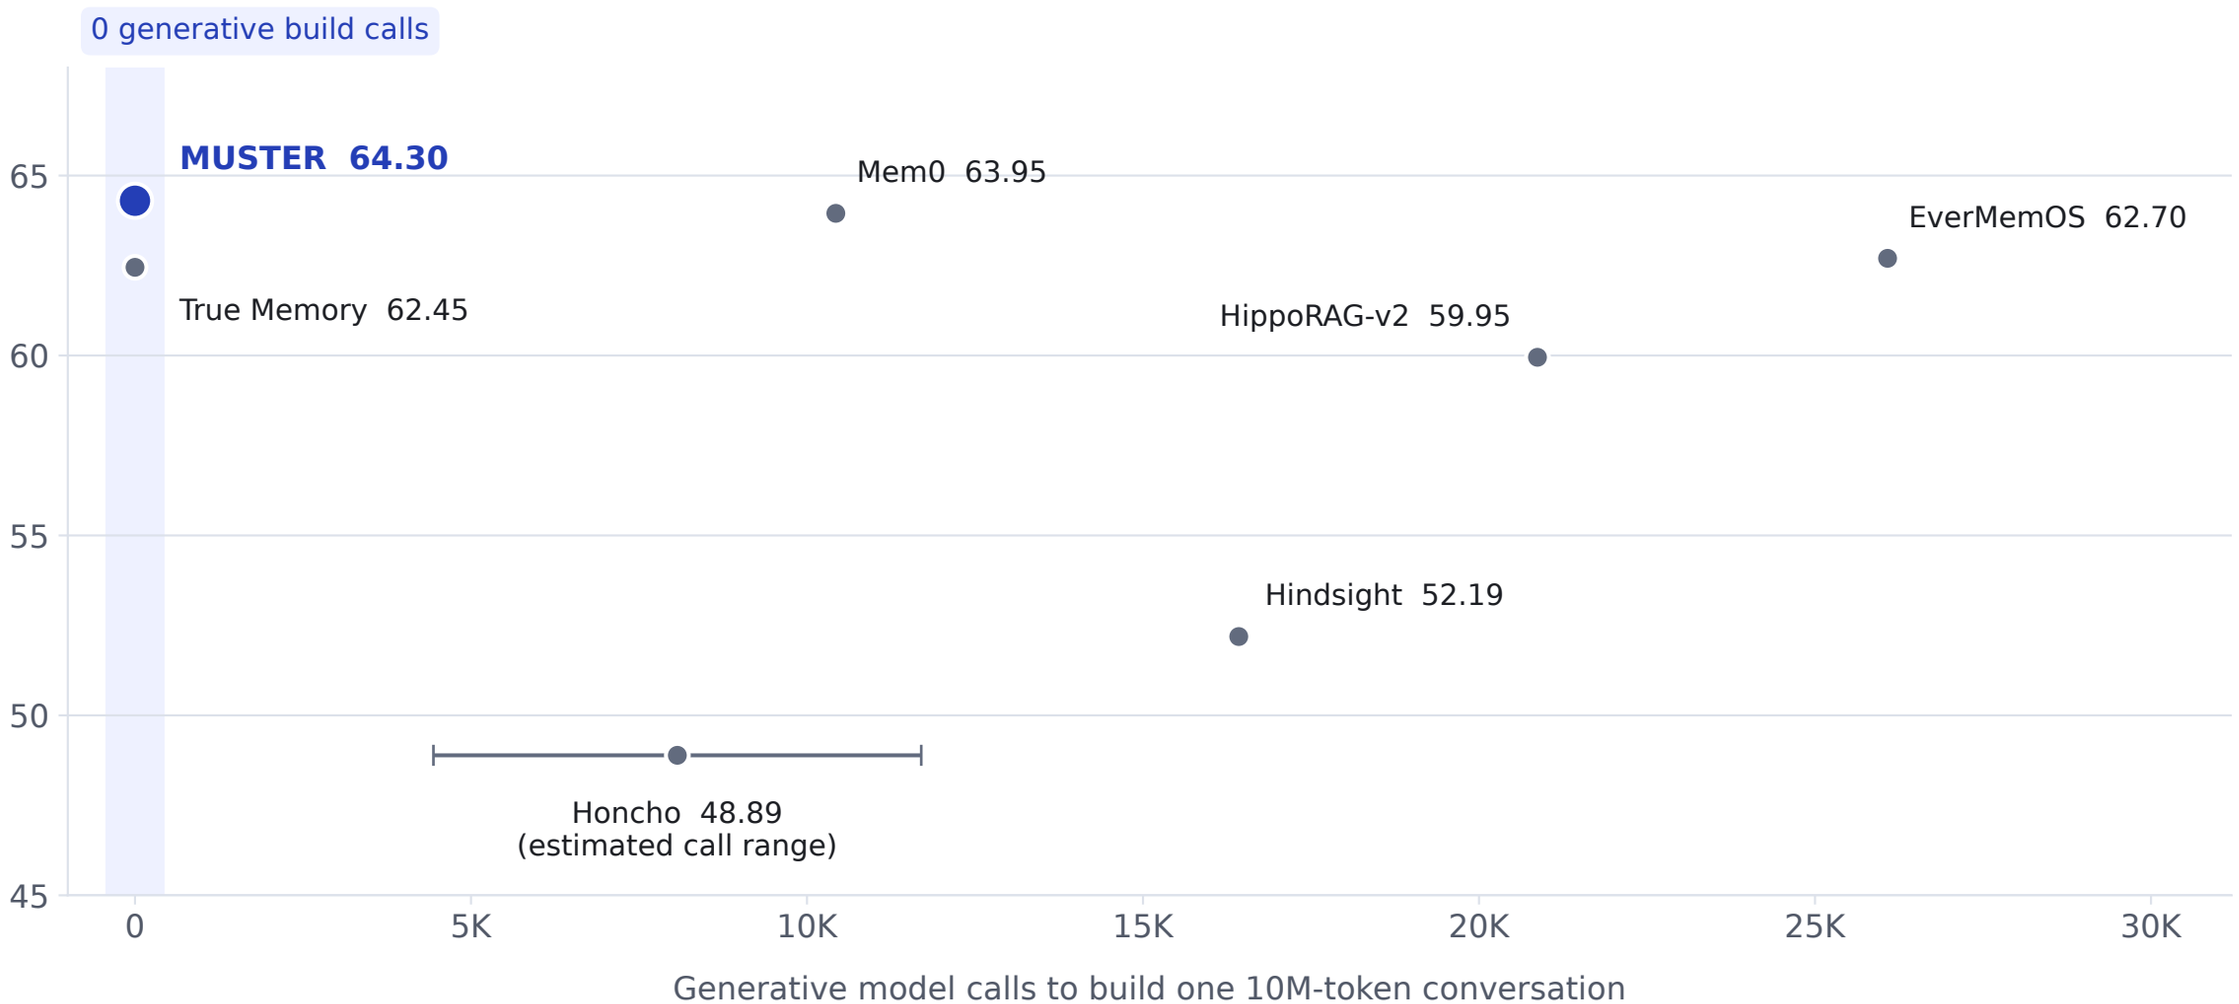

MUSTER: 77 s to build conversation 0 on one RTX PRO 6000 (fp16).

Zero generative calls does not mean zero embedding or reranking work; True Memory also uses zero.

Contexts and memory writers follow the paper's reported settings; this is not an equal-context comparison.

EverMemOS: 26,080 calls per 10M-token conversation (Table 15; rounded per-conversation mean).

Sources: MUSTER Table 7, BEAM results and build-cost appendix. Scope details and writers are in the caption / source ledger.